

# Advanced Message-Passing Programming

---

Parallel Filesystems and Lustre



# Reusing this material



This work is licensed under a Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International License.

<http://creativecommons.org/licenses/by-nc-sa/4.0/>

This means you are free to copy and redistribute the material and adapt and build on the material under the following terms: You must give appropriate credit, provide a link to the license and indicate if changes were made. If you adapt or build on the material you must distribute your work under the same license as the original.

Note that this presentation contains images owned by others. Please seek their permission before reusing these images.


# Overview

- Lecture will cover
  - Parallel Filesystems
  - Lustre Filesystem
  - Striping
  - Simple Lustre commands
  - Bottlenecks

# Parallel File Systems

- Parallel computer
  - constructed of many standard processors, each not particularly fast
  - performance comes from using many processors at once
  - requires manual distribution of data and calculation across processors
- Parallel file systems
  - constructed from many standard disks, each not particularly fast
  - performance comes from reading / writing to many disks at once
  - requires many *clients* to read / write to different disks
    - each *node* appears as a separate IO client
  - data from a single file can be *striped* across many disks
- Must appear as a single file system to user
  - typically have a single *MetaData* Server (MDS)

# Parallel File Systems: Lustre



# ARCHER's (not ARCHER2) Cray Sonexion Storage



## *SSU: Scalable Storage Unit*

2 x OSSs and 8 x OSTs (Object Storage Targets)

- Contains Storage controller, Lustre server, disk controller and RAID engine
- Each unit is 2 OSSs each with 4 OSTs of 10 (8+2) disks in a RAID6 array



**Multiple SSUs are combined to form storage racks**

# Terminology

- Lustre has many different levels and virtualisations
  - e.g. one Object Storage Server has multiple Object Storage Targets
  - a single OST has many physical disks in a RAID array
- I will refer to the following parts of Lustre
  - Meta Data Server (MDS)
    - the database that contains information on, e.g., where a file is stored
  - Object Storage Target
    - the physical device that stores your data
    - I may also call this a “disk” (although it contains multiple hard drives)
- The MDS and the OSTs are what a user interacts with

# ARCHER2 hardware

- Three separate /work filesystems (work1, work2 & work3)
  - each has 12 OSTs and one MDS
  - consortia assigned to different partitions to share the load
  - multiple filesystems means the MDS is less likely to be overloaded
- One filesystem with Solid State (SSD) not spinning disks
  - still to be configured
  - expect better latency, e.g. good for small I/O transactions
- Each has around 3.3 PiB of storage
  - total of 13.2 PiB (which is 14.5 PB !)





# Default Configuration

- By default, each file is stored on a single OST
  - assigned when the file is created
  - distributed across all available OSTs to balance the load
  - each OST is actually a separate Linux filesystem
- This is called an “unstriped” file
- Reading and writing multiple files from multiple nodes can benefit from multiple OSTs
- Access to a single file will not benefit from the parallel nature of the filesystem

# Lustre data striping

Parallel performance comes from striping single files over multiple OSTs



OS/file-system  
divides the file into  
stripes *if requested*  
*by the user*

Stripes read/written to/from  
their assigned OST

Single logical user  
file e.g.  
/work/q01/q01/user  
/bigfile.dat

# Striping

- Allow multiple IO processes to access same file
  - increases bandwidth as you are accessing multiple OSTs
- Typically optimised for bandwidth, not for latency
  - e.g. reading/writing small amounts of data is very inefficient
- This is called striping
  - striping of a file is fixed when it is created, under control of the user
  - fundamental parameters are the number of stripes and stripe size
- For example, if a file is created with 4 stripes
  - Lustre assigns four OSTs: OST1, OST2, OST3, OST4
  - first MiB is stored on OST1, second on OST2, third on OST3, fourth on OST4, fifth on OST1, sixth on OST2, ....
  - i.e. round-robin with default stripe size of 1MiB

# Lustre commands

- To set the striping on a directory or file

```
lfs setstripe -c nstripe <dir/file>
```

- nstripe = -1 is full striping (12 on each of ARCHER2's 3 filesystems)

- Stripe size: 

```
lfs setstripe -s 4m <dir/file>
```


- Does **not** alter striping for existing files: that requires a copy
- I always use setstripe on directories
  - all files subsequently created in directory will have the same striping
- To enquire: 

```
lfs getstripe <dir/file>
```

# Parallel IO to a striped file

- Very complicated in practice!
  - where in the file does the local data need to be written?
  - which OSTs are the stripes located on?
  - are there write conflicts coming from different processes?
- Need to use a parallel IO library

# Lustre: where are the bottlenecks?



# Benchio benchmark

- Obvious questions:
  - does the MDS become overloaded for large numbers of files?
  - what is the maximum performance of a single OST?
  - can a single process saturate an OST?
  - can a single node saturate an OST?
    - or is the network the limiting factor
  - how well do different IO libraries work with Lustre?
  - what are the best stripe count (and size) settings?
  - ....
- I wrote a simple benchmark to help investigate Lustre performance characteristics and bottlenecks
  - we will use benchio for the practical examples

# Summary

- A Lustre filesystem has multiple OSTs
  - I think of these as being multiple disks
- By default on ARCHER2, each file stored on a single OST
  - i.e. an unstriped file with a stripe count of 1
  - increased performance for multiple files
    - a single user writing many files
    - multiple users each writing a single file
- Improving performance for a single file requires striping
  - fully under control of the user
  - expect parallel IO libraries to take advantage of striping